



# Thinking in a Highly Concurrent, Mostly-functional Language

Elixir Conf

Austin, July 26<sup>th</sup> 2014

**Francesco Cesarini**  
Founder & Technical Director

@francescoC  
[francesco@erlang-solutions.com](mailto:francesco@erlang-solutions.com)



Erlang Training and Consulting Ltd

# Thinking in a Highly Concurrent, Mostly-functional Language

QCON London, March 12<sup>th</sup>, 2009

**Francesco Cesarini**  
[francesco@erlang-consulting.com](mailto:francesco@erlang-consulting.com)

```
counter_loop(Count) ->
    receive
        increment ->
            counter_loop(Count + 1);
        {count, To} ->
            To ! {count, Count},
            counter_loop(Count)
    end.
```

# Erlang



**After you've opened the top of your head,  
reached in and turned your brain inside out,  
this starts to look like a natural way to count  
integers. And Erlang does require some fairly  
serious mental readjustment.**

**However... having spent some time playing with  
this, I tell you...**

Tim Bray, Director of Web Technologies - Sun Microsystems

**... If somebody came to me and wanted to pay me a lot of money to build a large scale message handling system that really had to be up all the time, could never afford to go down for years at the time, I would unhesitatingly choose Erlang to build it in.**

Tim Bray, Director of Web Technologies - Sun Microsystems

# Syntax

# Concurrency

# Erlang Highlights: Concurrency

## *Creating a new process using spawn*

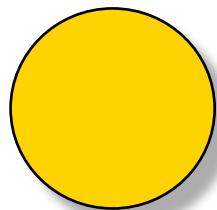
```
-module(ex3).  
-export([activity/3]).  
  
activity(Name,Pos,Size) ->  
.....
```



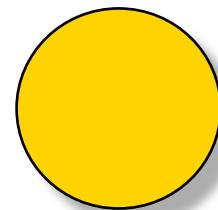
*Pid = spawn(ex3, activity, [Joe, 75, 1024])*

# Erlang Highlights: Concurrency

*Processes communicate by asynchronous message passing*



Pid ! {data,12,13}



```
receive
  {start} -> .....
  {stop} -> .....
  {data,X,Y} -> .....
end
```

# Products: AXD301 Switch - 1996

A Telephony-Class, scalable (10 - 160 GBps) ATM switch

Designed from scratch in less than 3 years

AXD 301 Success factors:

- Competent organisation and people
- Efficient process
- Excellent technology (e.g. Erlang/OTP)



# Products: AXD301 Switch - 1996

Erlang: ca 1.5 million lines of code

- Nearly all the complex control logic
- Operation & Maintenance
- Web server and runtime HTML/  
JavaScript generation

C/C++: ca 500k lines of code

- Third party software
- Low-level protocol drivers
- Device drivers

Java: ca 13k lines of code

- Operator GUI applets



# Concurrency Modeling

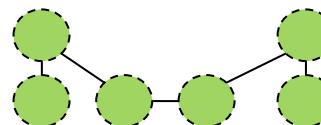
Model for the natural concurrency in your problem

In the old days, processes were a critical resource

- Rationing processes led to complex and unmanageable code

Nowadays, processes are very cheap: if you need a process - create one!

Example: AXD301 process model



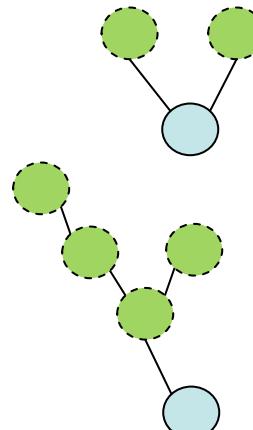
1<sup>st</sup> prototype:  
6 processes/call



2 processes/call



1 process/all calls



2 processes/  
call transaction

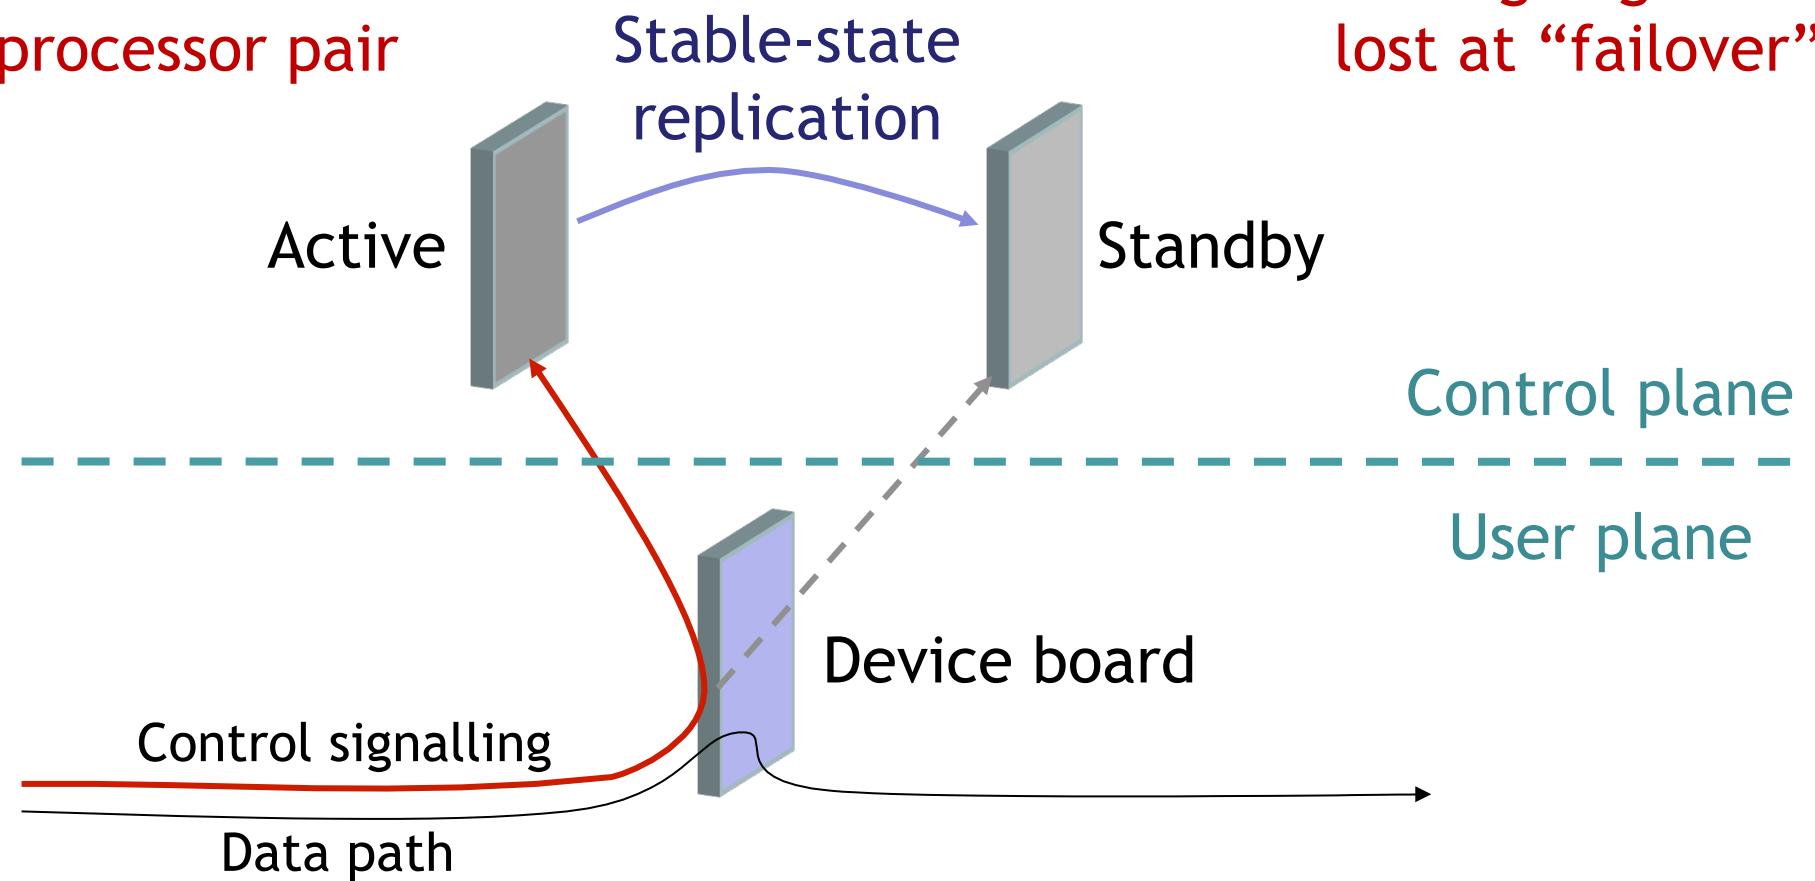


4-5 processes/  
call transaction

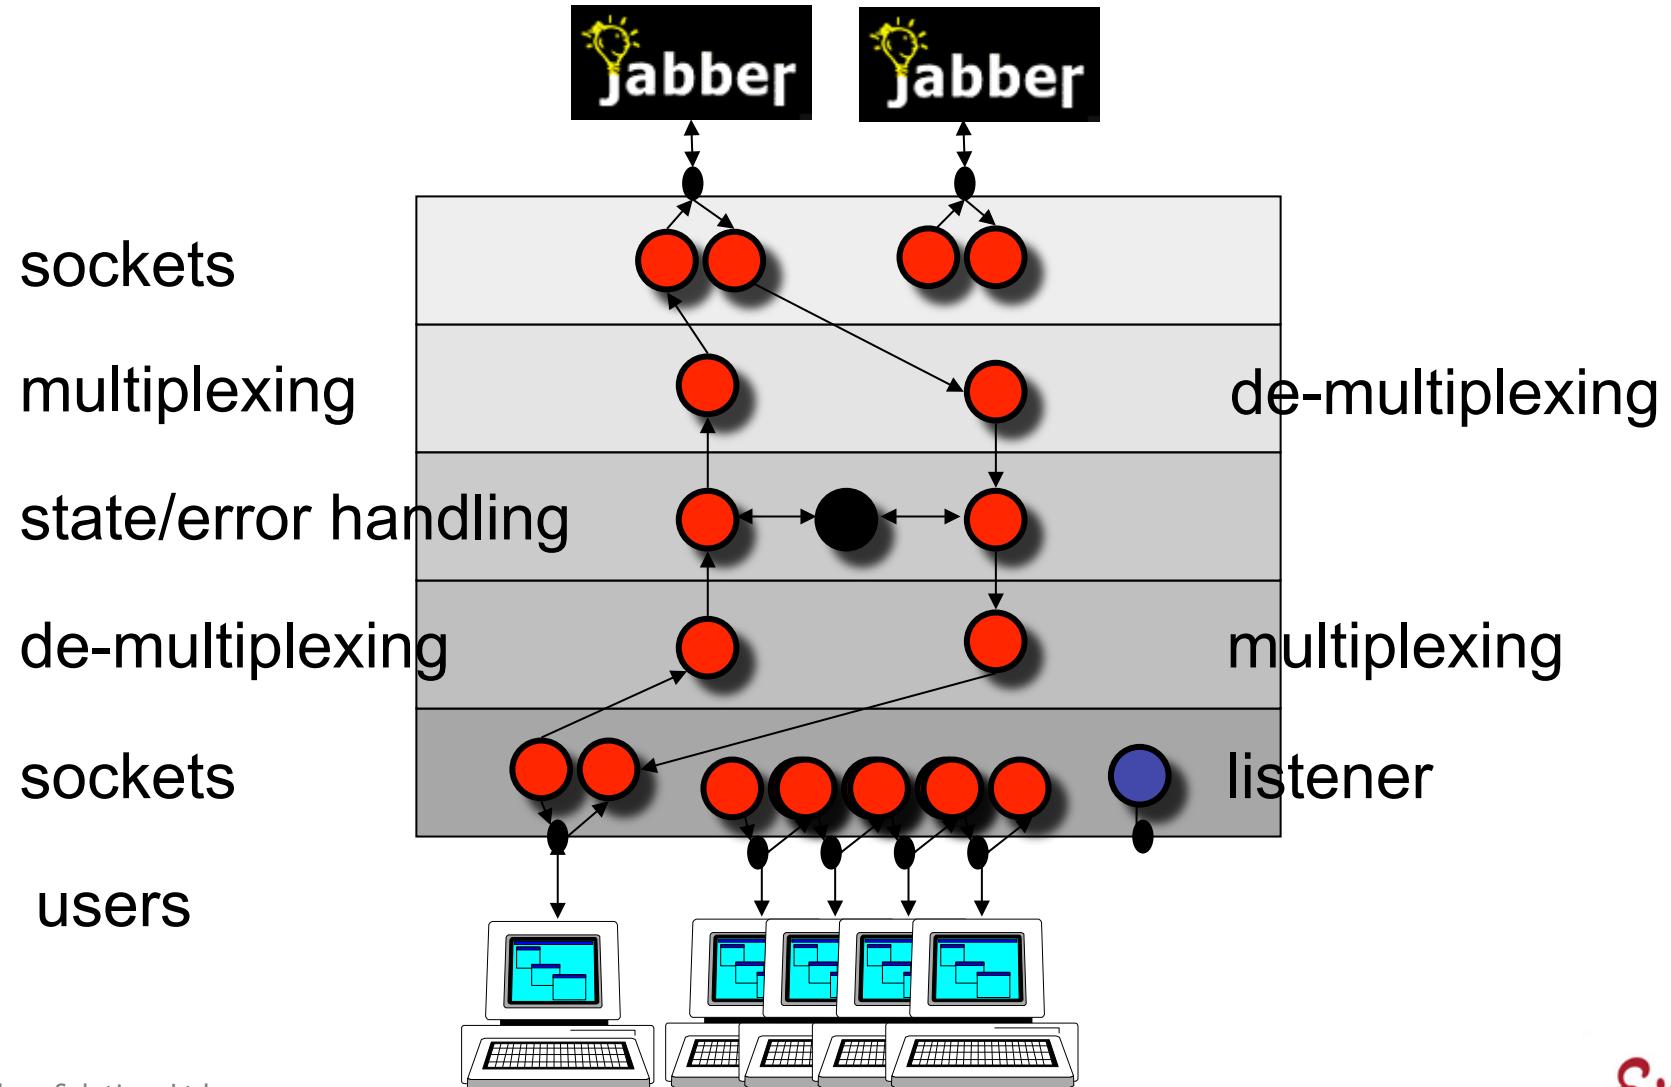
# 1+1 Redundancy - Good ol' Telecoms

~ 35 000 calls  
per processor pair

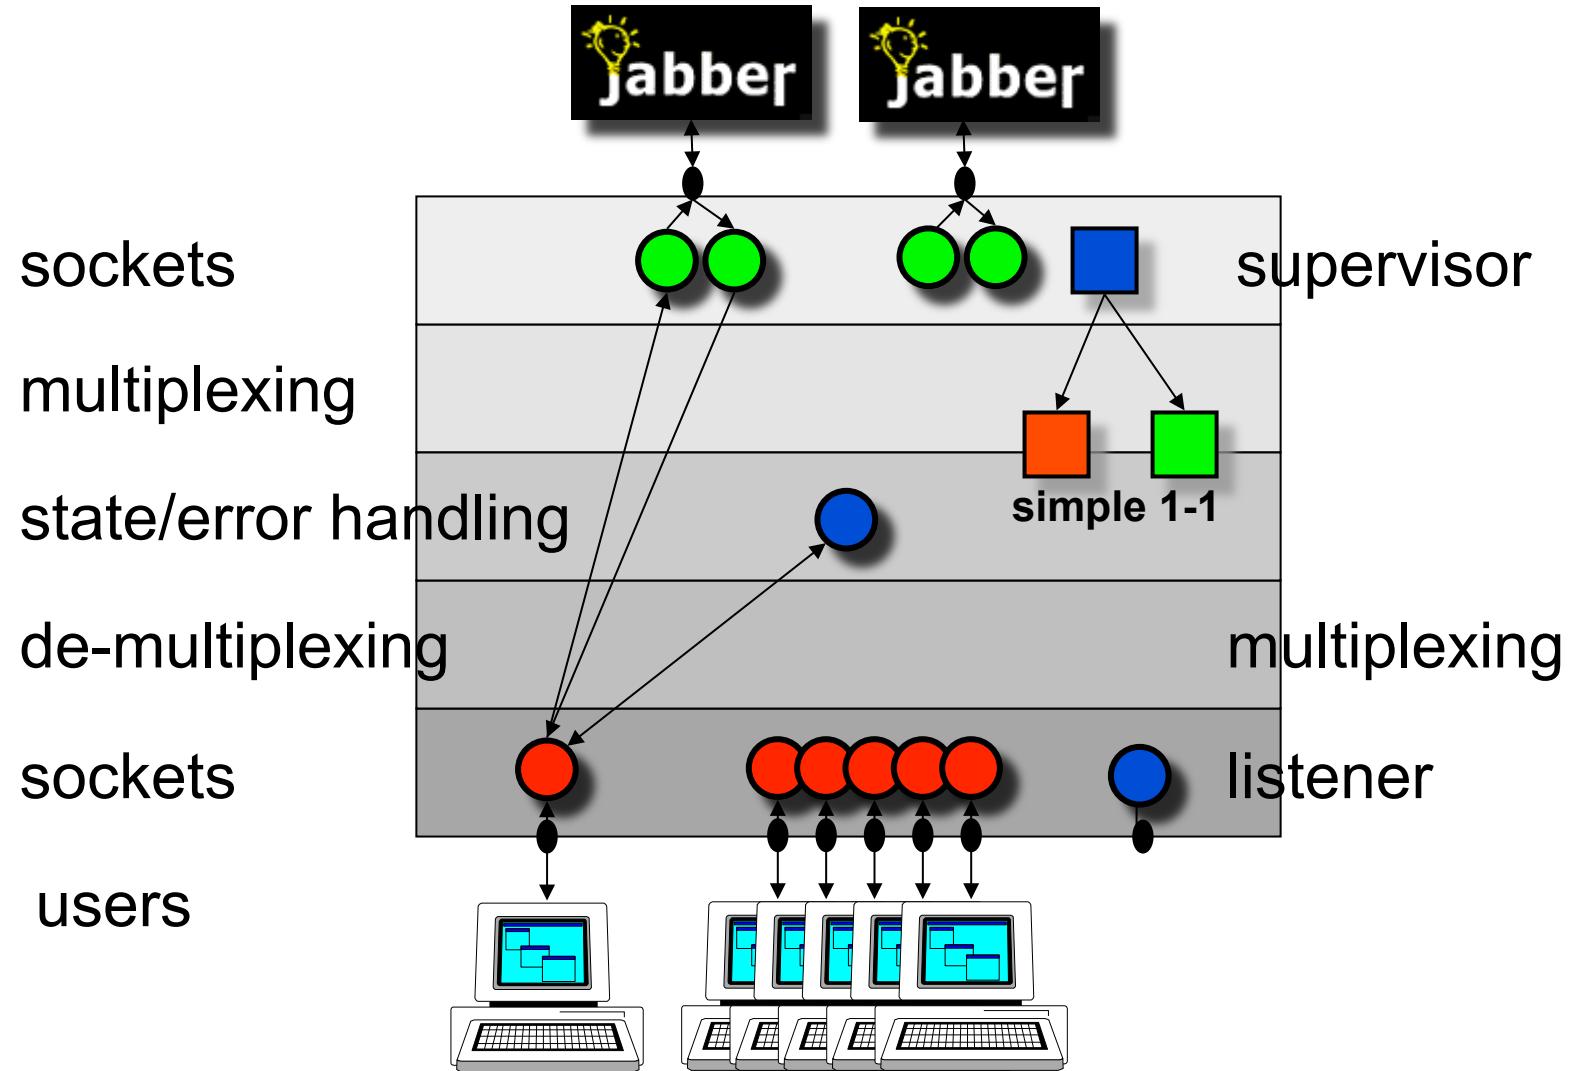
No ongoing sessions  
lost at “failover”



# First IM Proxy Prototype - 2000



# First IM Proxy Prototype - 2000



# Products: EjabberD IM Server - 2002

A distributed XMPP server

Started as an Open Source  
Project by *Alexey Shchepin*

Commercially Supported by  
Process-One (Paris)

- 40% of the XMPP IM market
- Used as a transport layer
- Manages 30,000 users / node



# Products: EjabberD IM Server - 2002

A distributed XMPP server

Started as an Open Source  
Project by *Alexey Shchepin*

Commercially Supported by  
Process-One (Paris)

- 40% of the XMPP IM market
- Used as a transport layer
- 2008, Manages 30,000 users / node

MongooseIM is a fork and rewrite

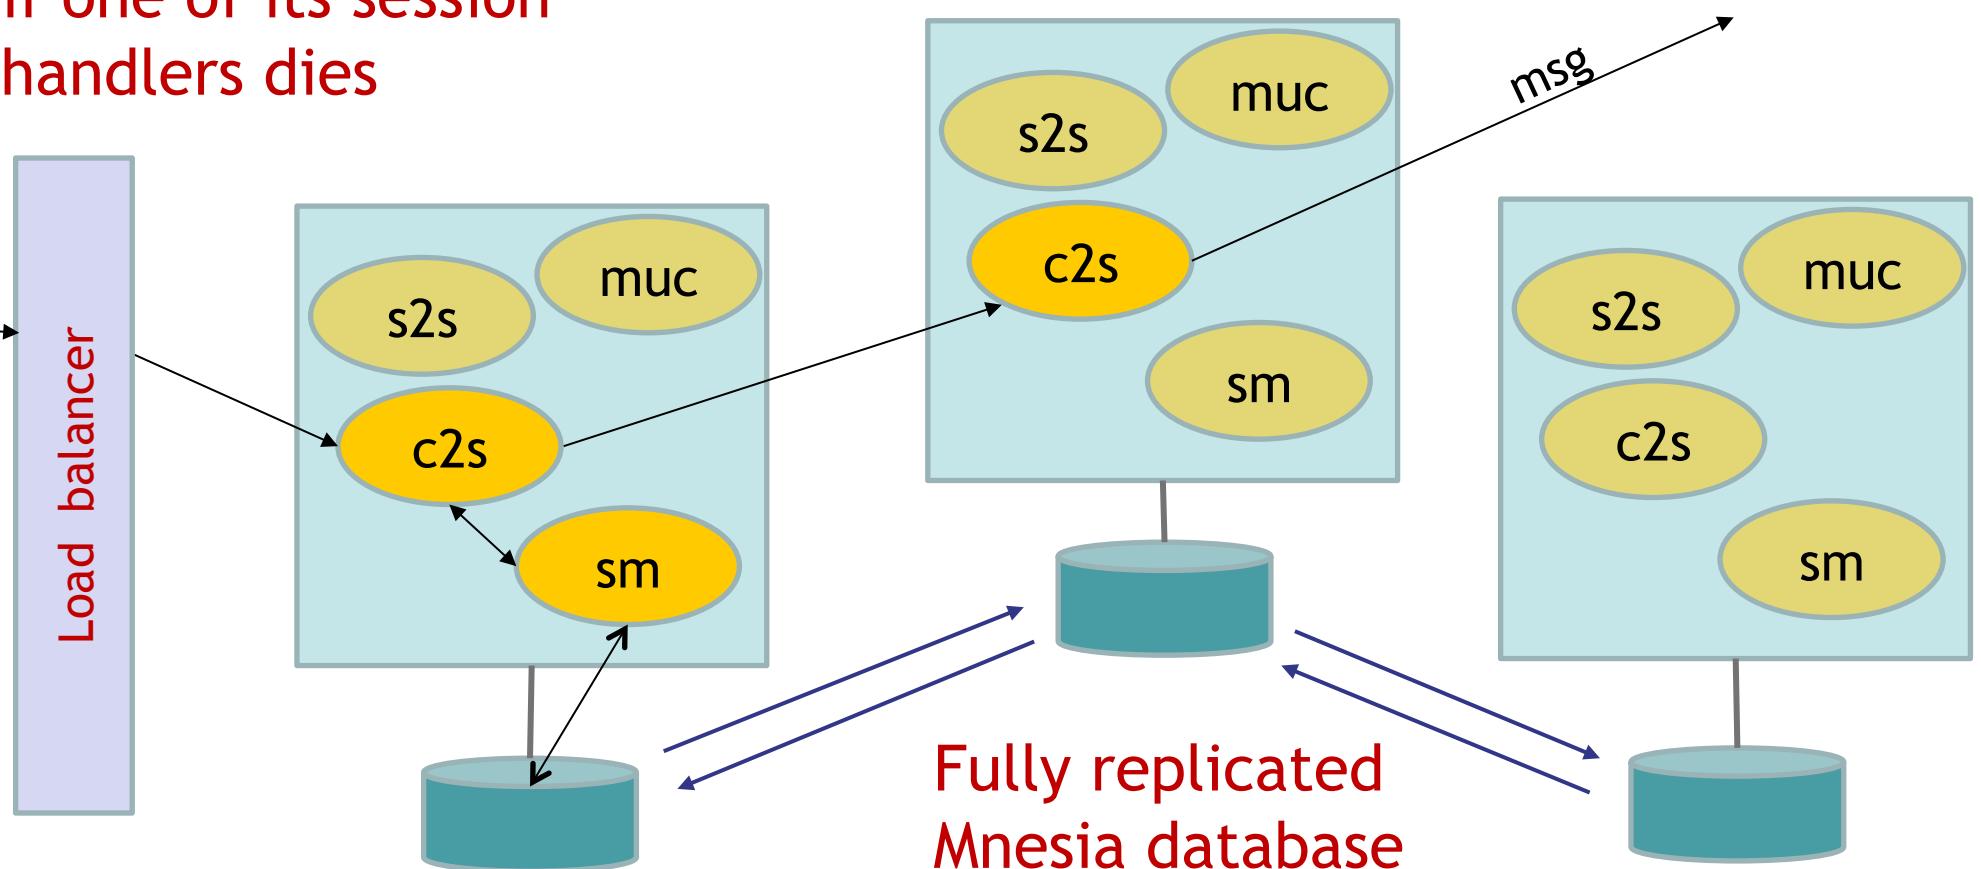
- Open Source, supported by Erlang Solutions
- Used for Messaging and Device Management
- 2014, reached 1 million users / node

© 2014 - Erlang Solutions Ltd

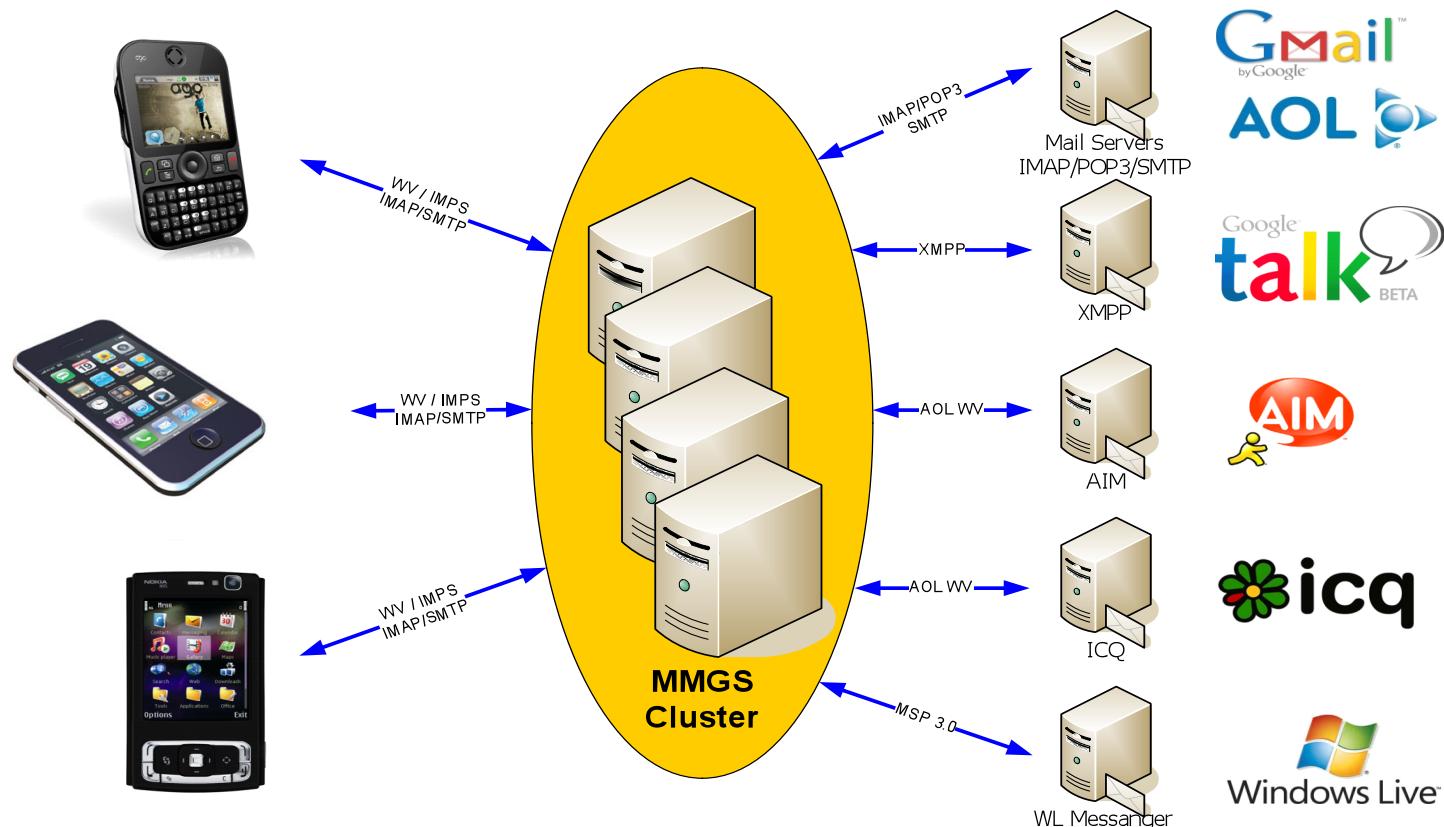


# Fully Replicated Cluster - Ejabberd 2002

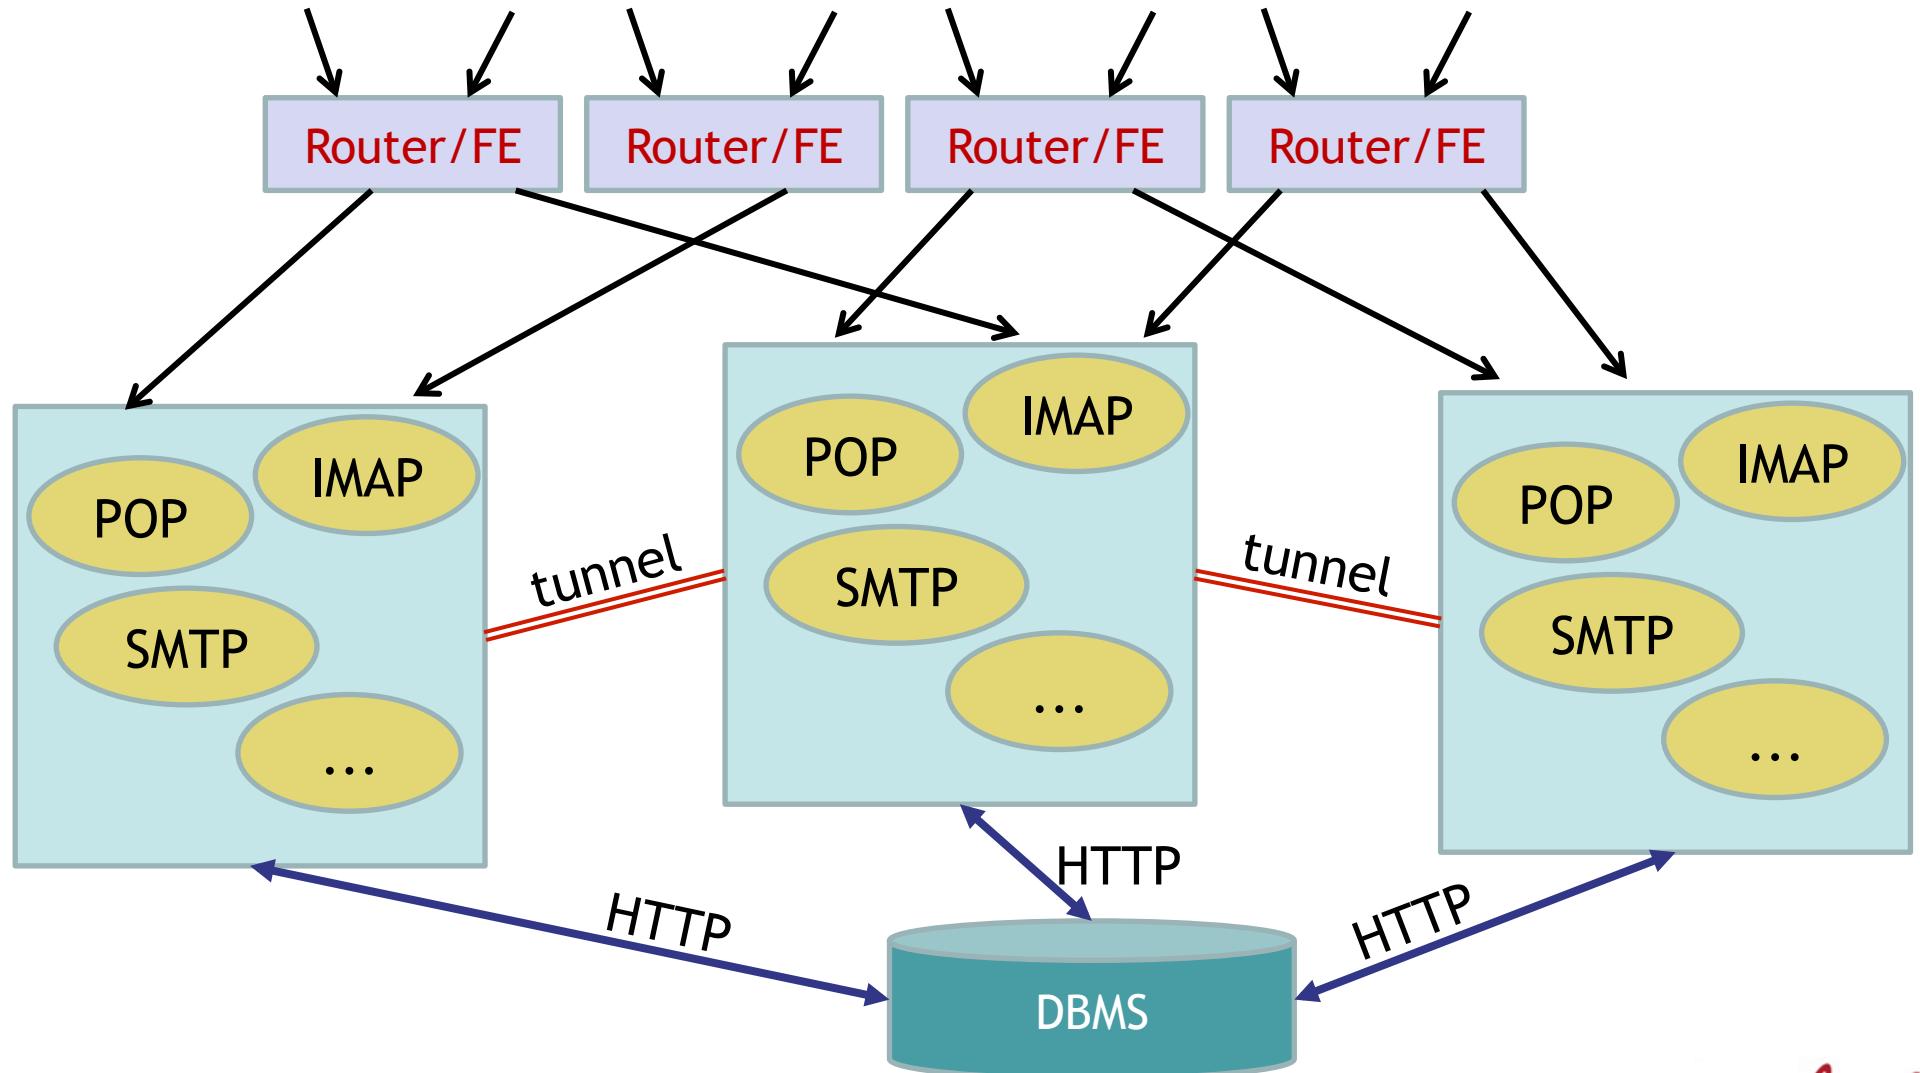
Client must re-connect  
if one of its session  
handlers dies



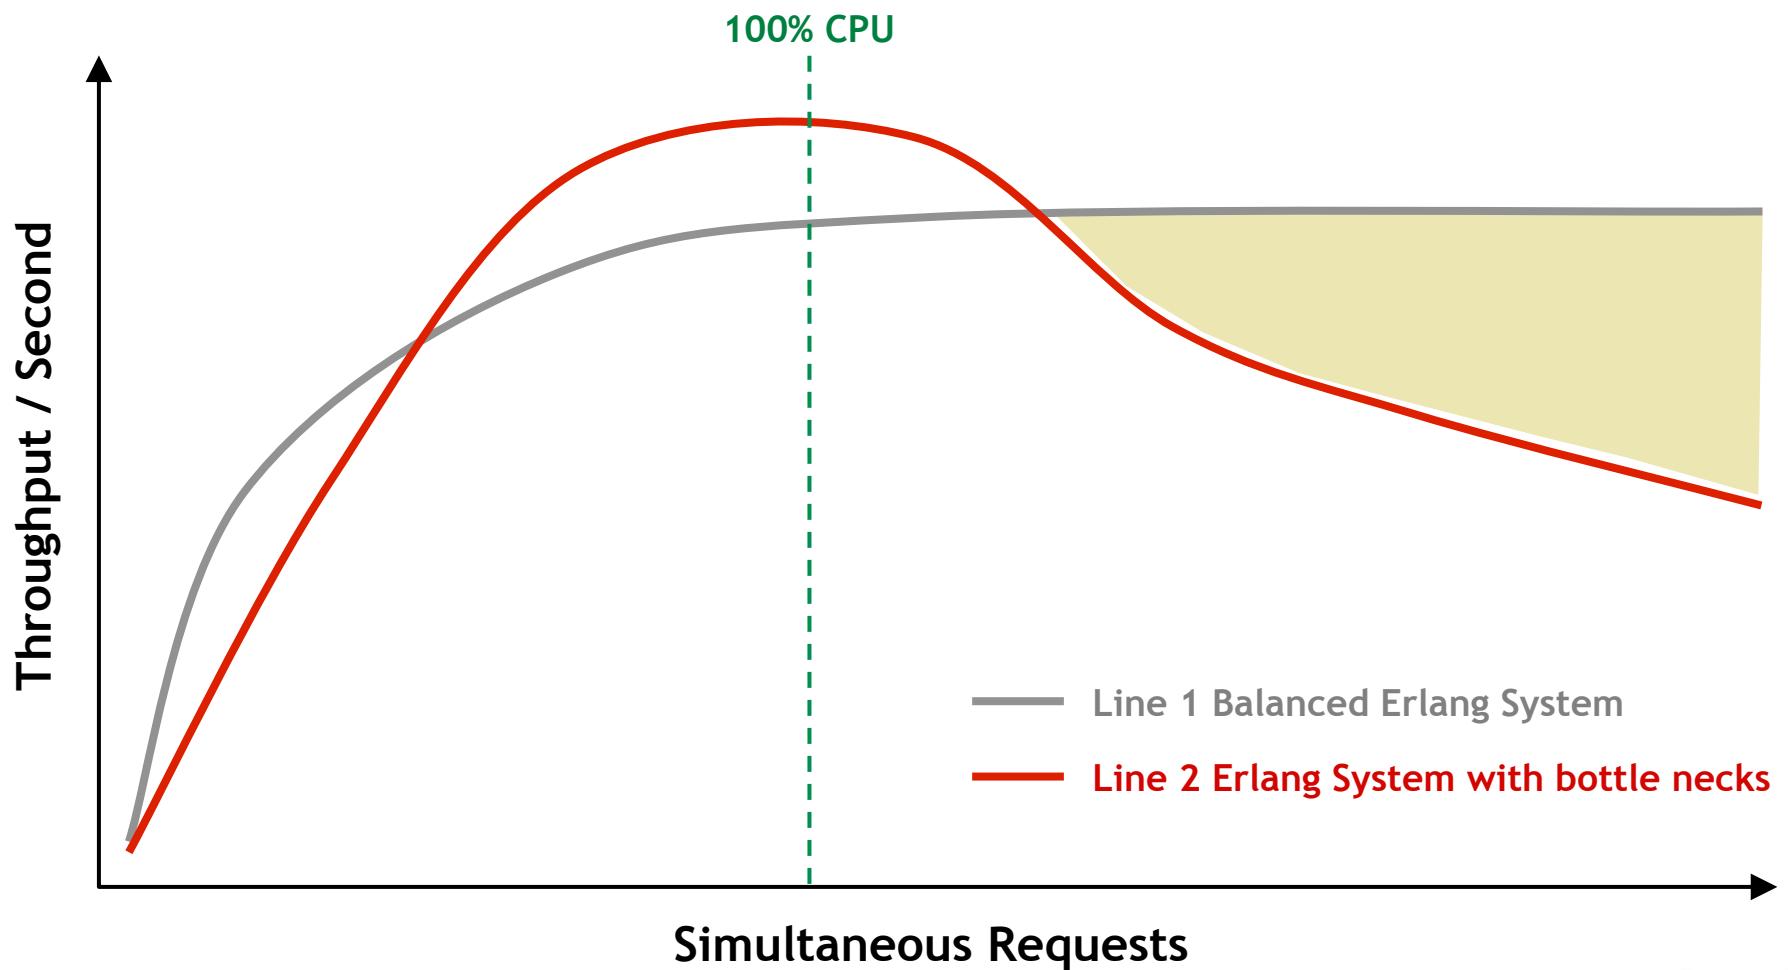
# Share-nothing Architecture - Messaging Gateway



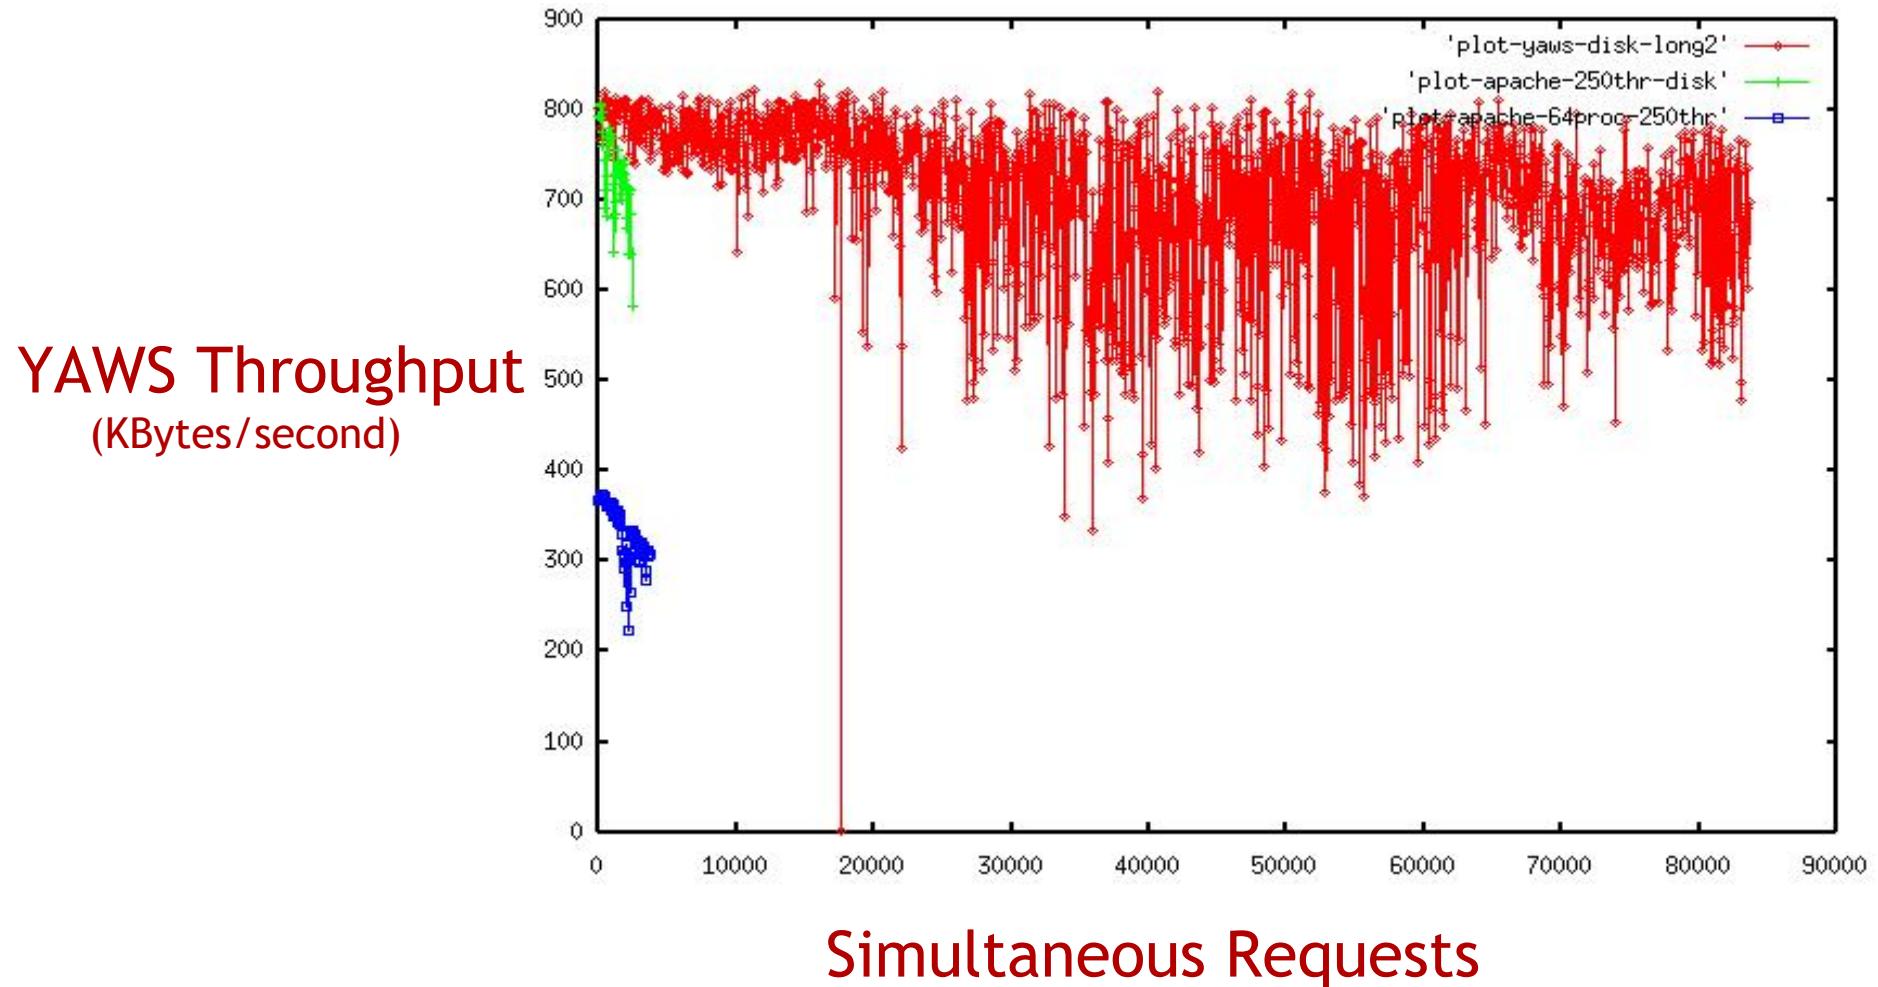
# Share-nothing Architecture - Messaging Gateway



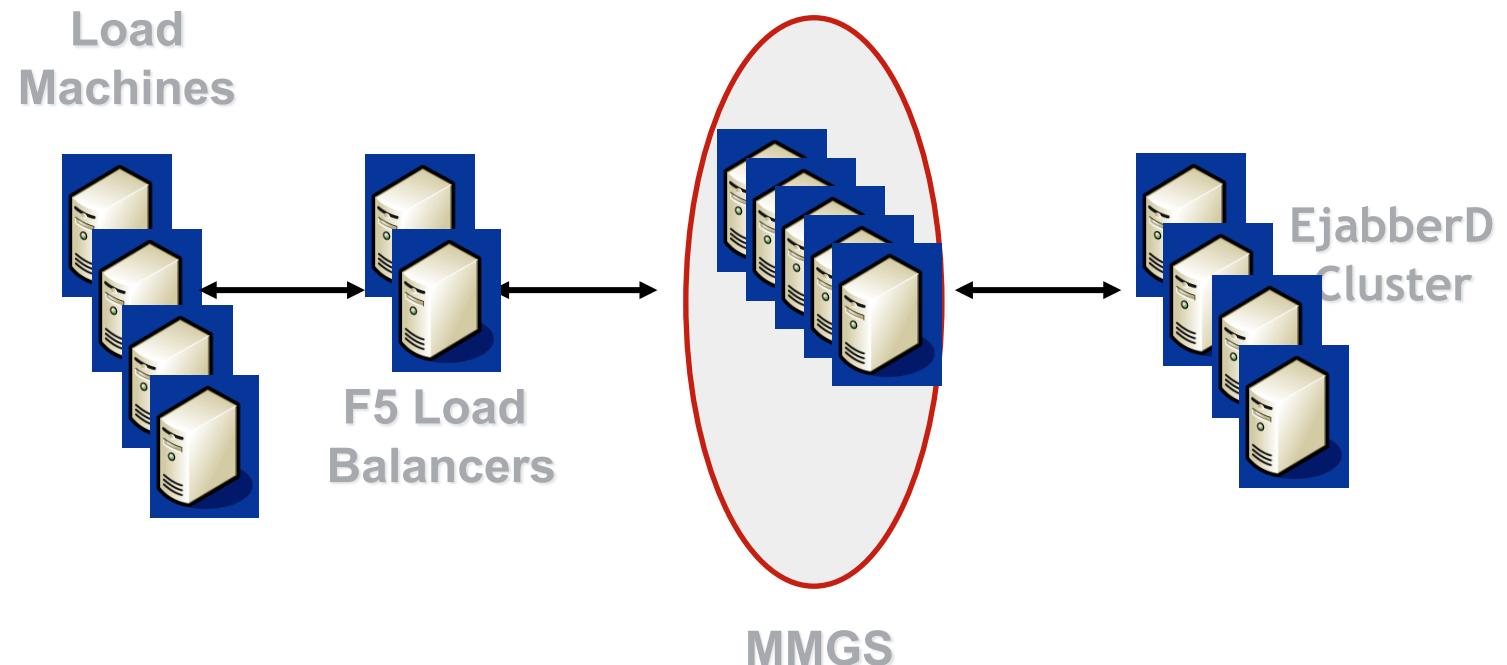
# Erlang Concurrency Under Stress - Pre-SMP



# Erlang Concurrency Under Stress - Pre-SMP



# Erlang Concurrency Under Stress - Post-SMP



# Stress Tests With SMP

I/O Starvation

TCP/IP Congestion

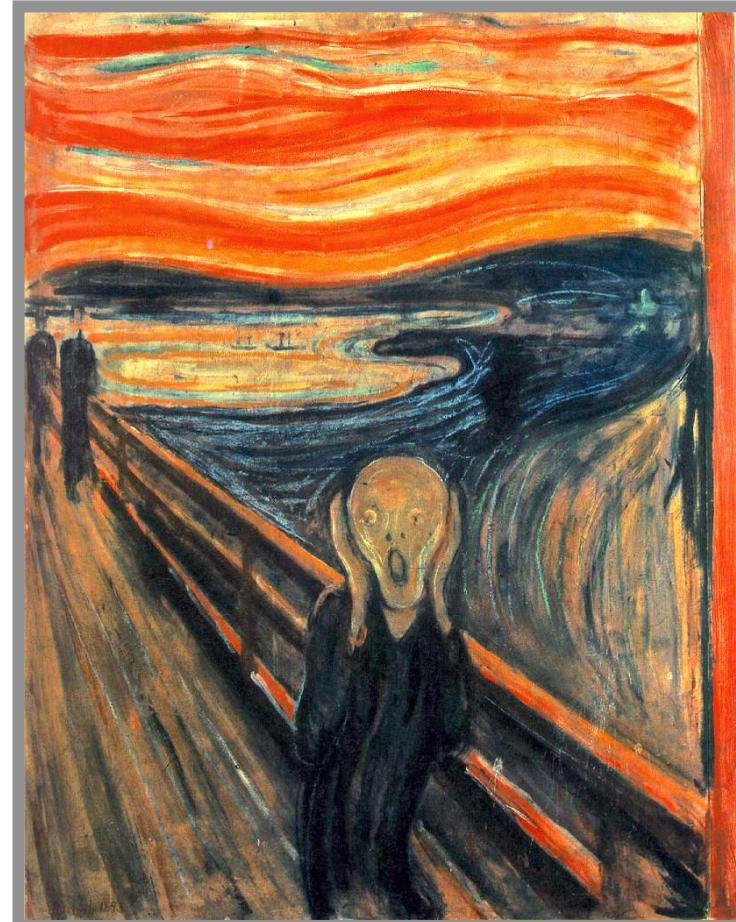
Memory Spikes

Timeout Fine-tuning

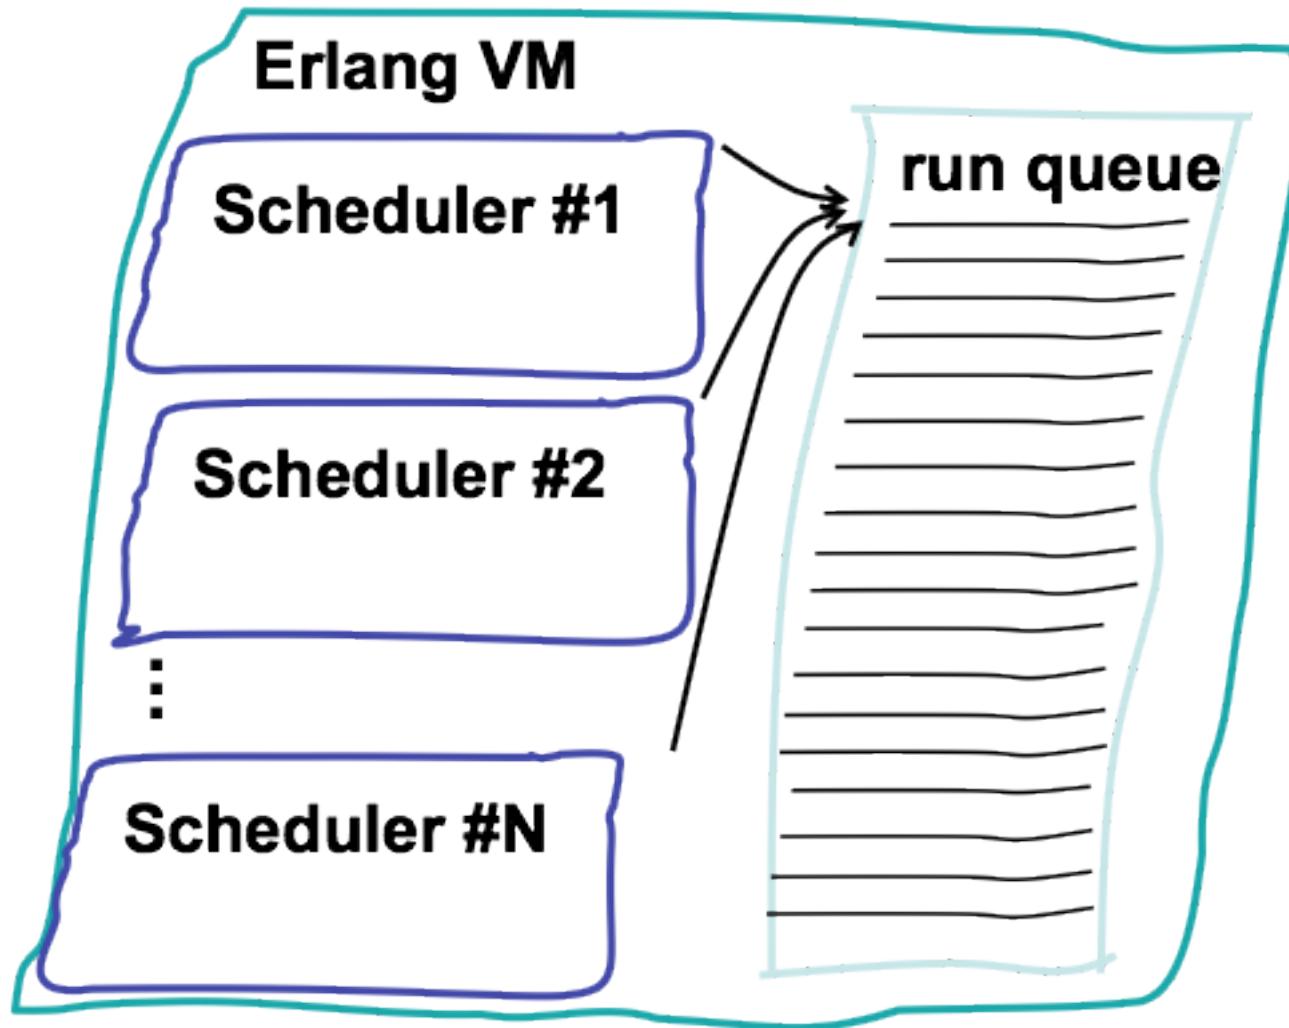
OS Limitations

ERTS Configuration Flags

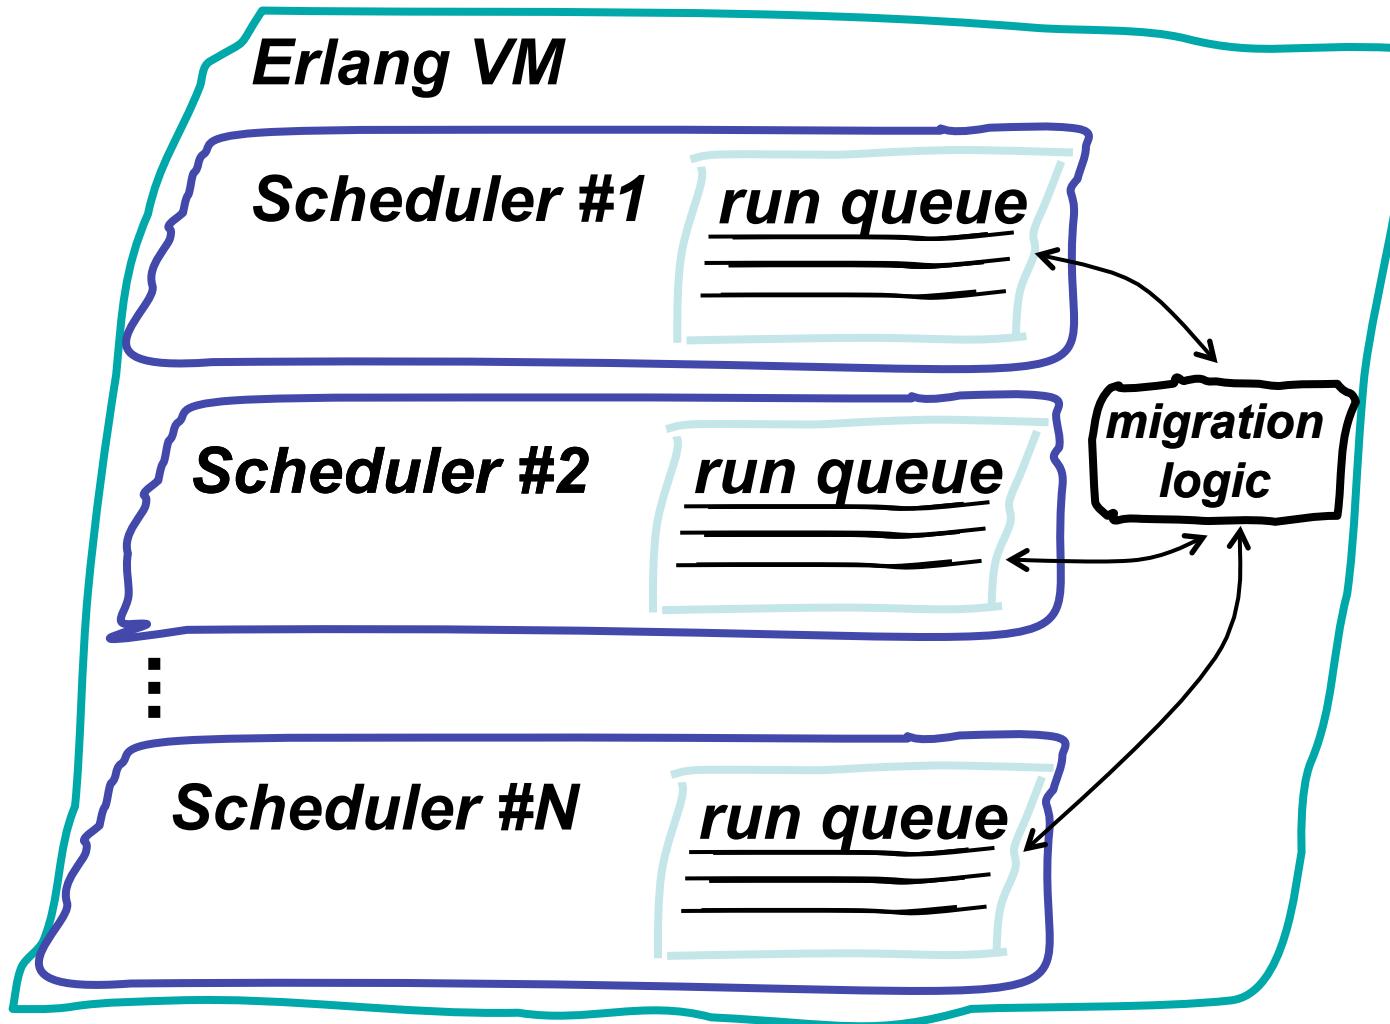
Shut down Audit Logs



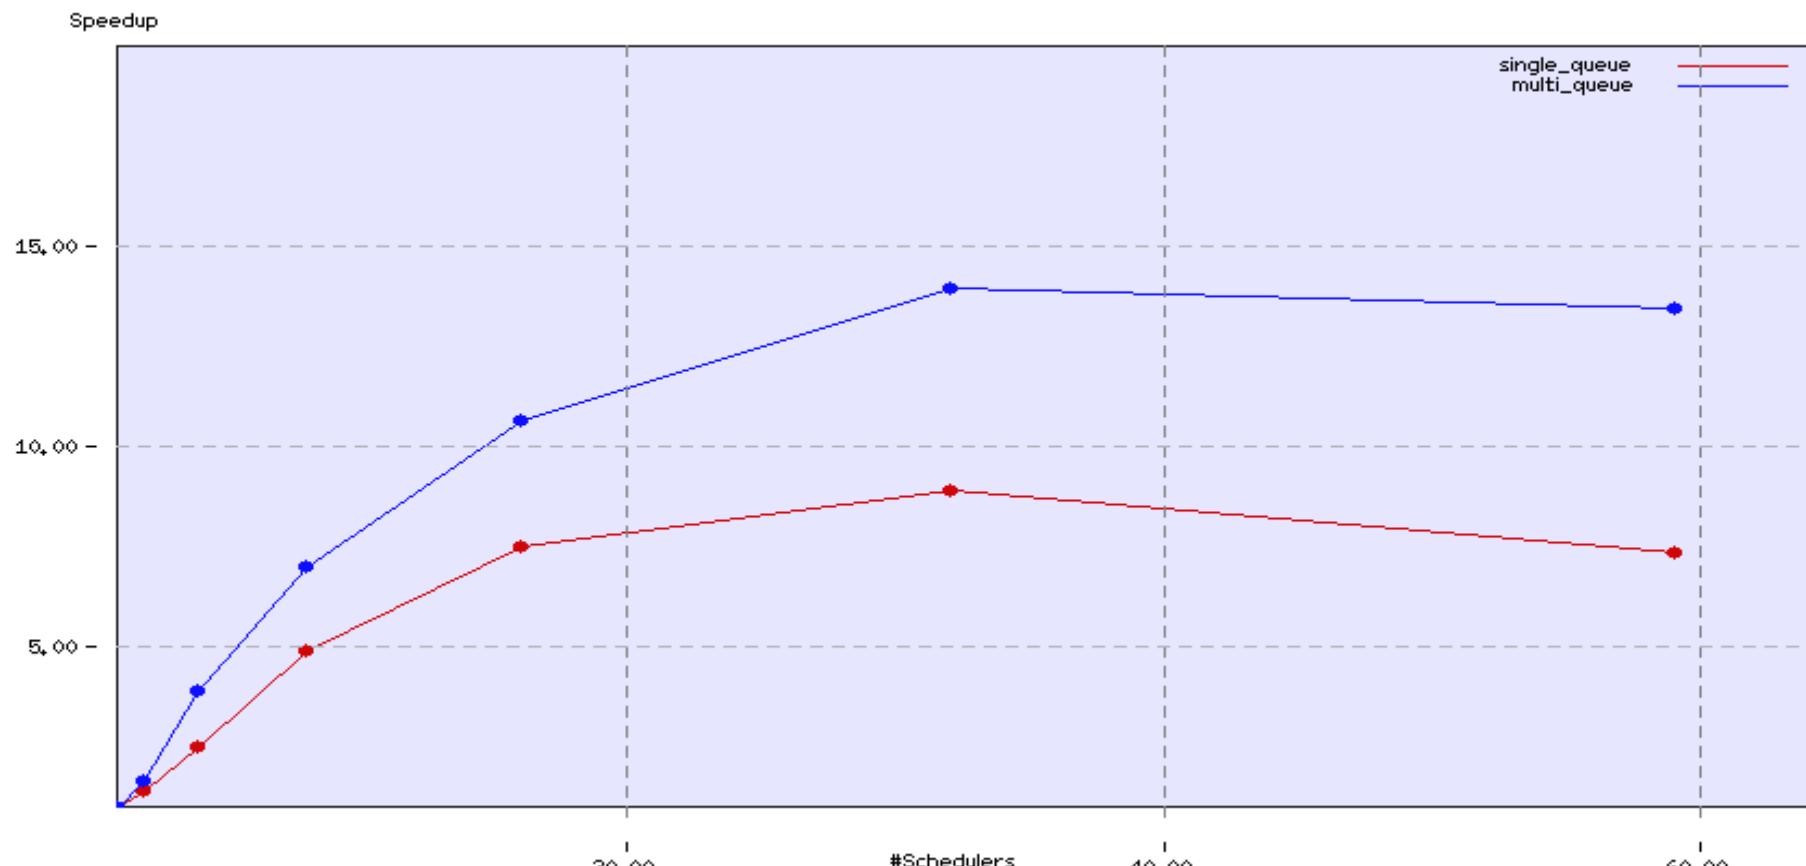
# SMP bottlenecks - pre 2008



# SMP bottlenecks - post 2008



# Big Bang Benchmark - post 2008

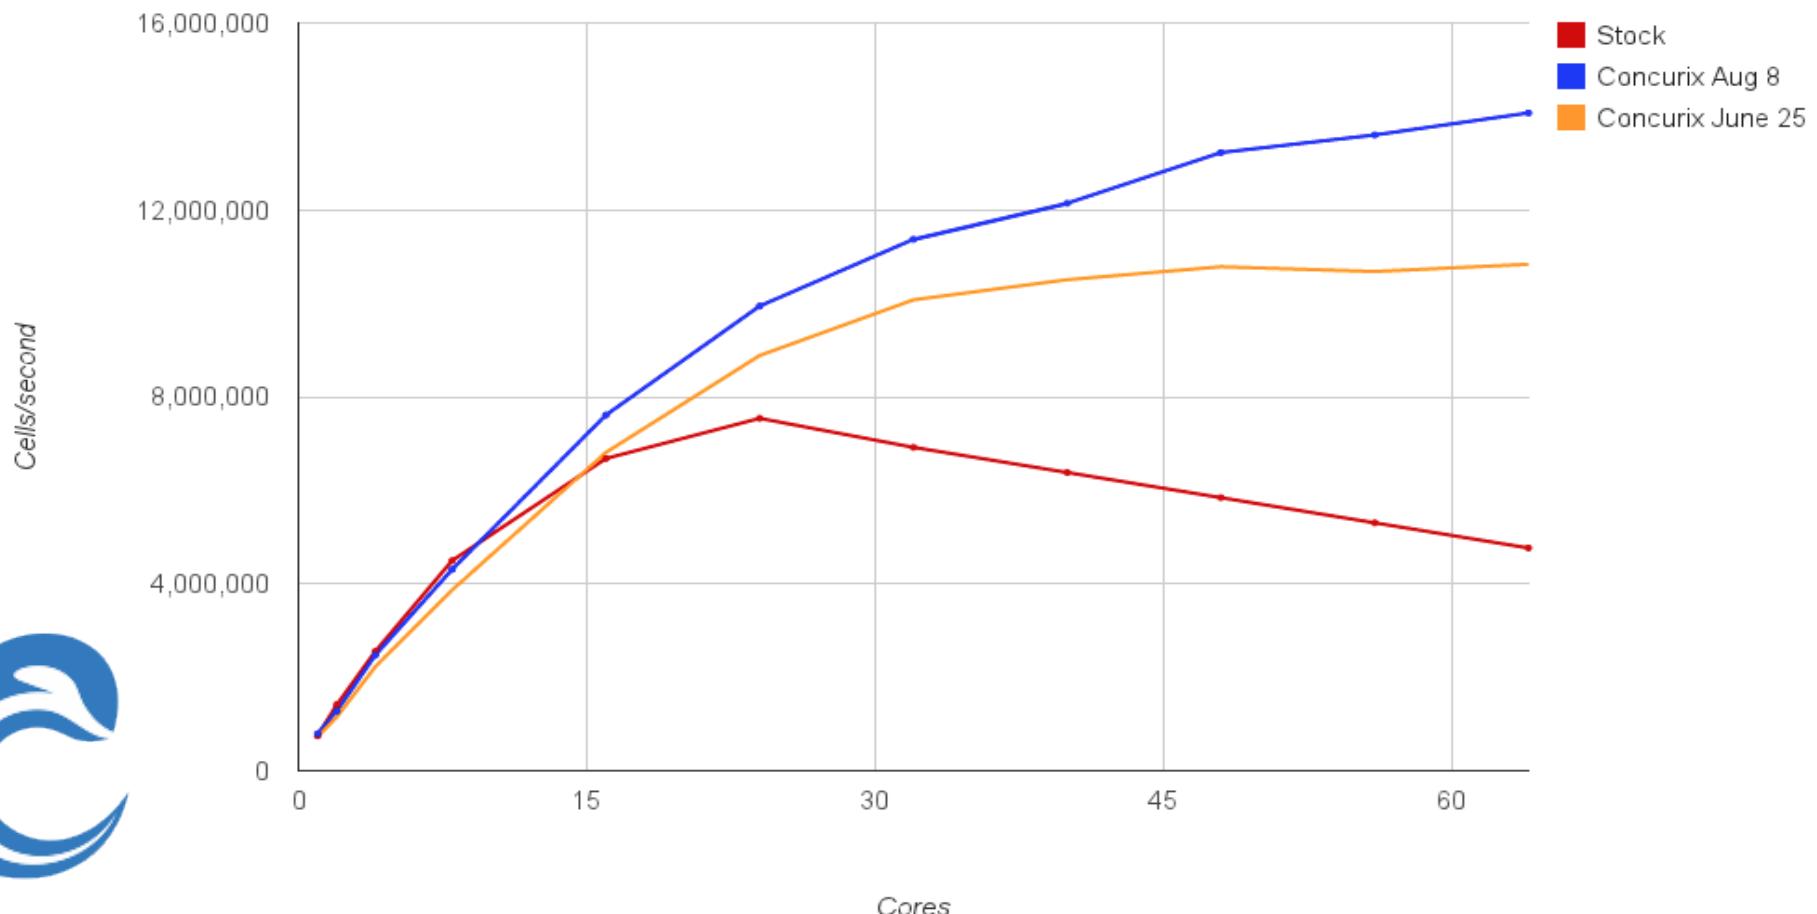


*Red: Single Queue, Blue: Multiple Run Queue on a Tilera TilePro64 (64 cores)*

# Mandelbrot- 2013



**Mandelbrot throughput**



# Now for the Bottlenecks



The screenshot shows the Chicago Boss website. At the top, there's a dark header with the text ".//Chicago BOSS" in white and orange. Below the header is a navigation bar with links for "API", "WIKI", "QUICK START", "DOWNLOAD V 0.8.4" (which is highlighted in orange), and "ABOUT". The main content area features a large, bold, black text: "Build your next website with Erlang — the world's most advanced networking platform." To the left of this text is a circular icon containing a white hourglass.

Build your next website with Erlang —  
the world's most advanced networking  
platform.



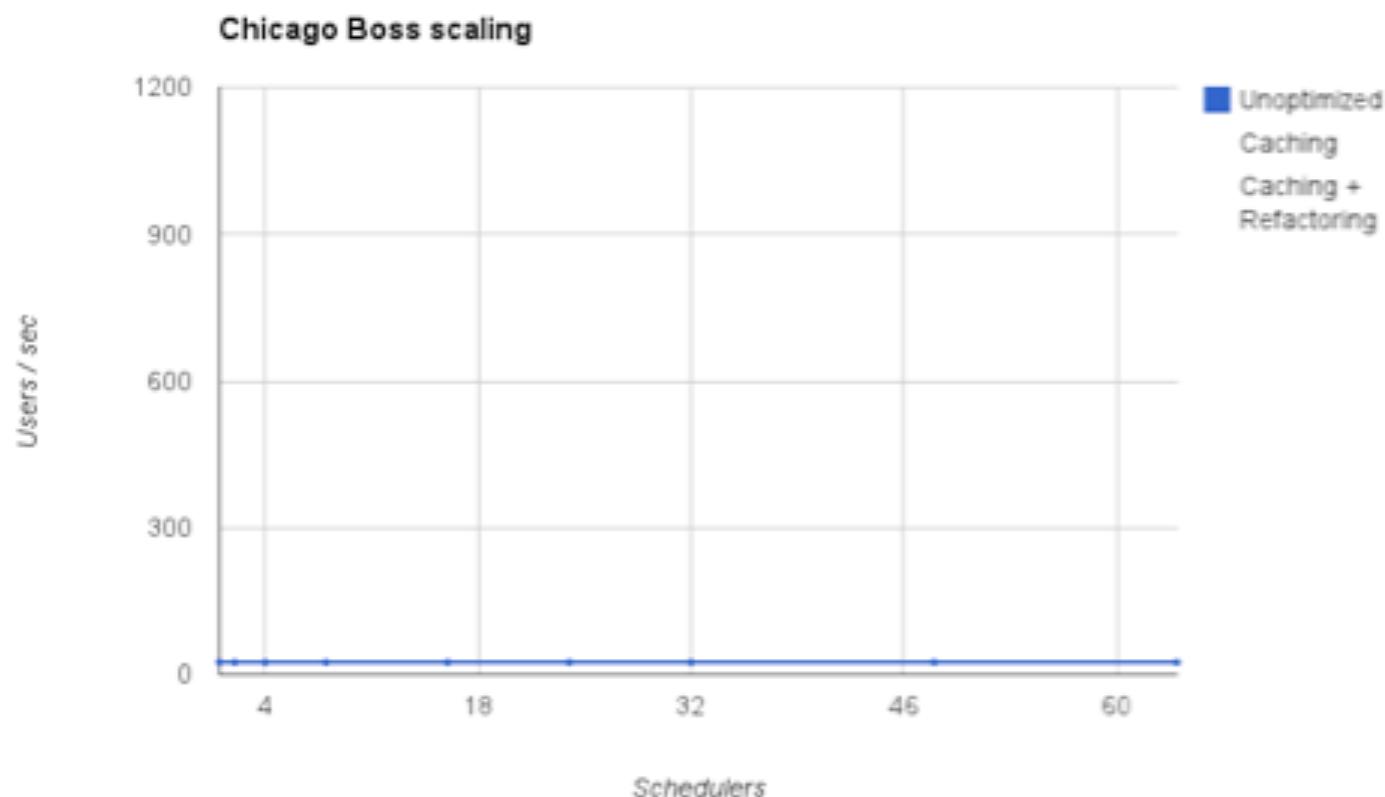
Do you pine for a simpler time when web pages loaded in under one second? **Chicago Boss** is the answer to slow server software: a Rails-like framework for Erlang that delivers web pages to your users as quickly and efficiently as possible.



© 2014 - Erlang Solutions Ltd

*Erlang*  
SOLUTIONS

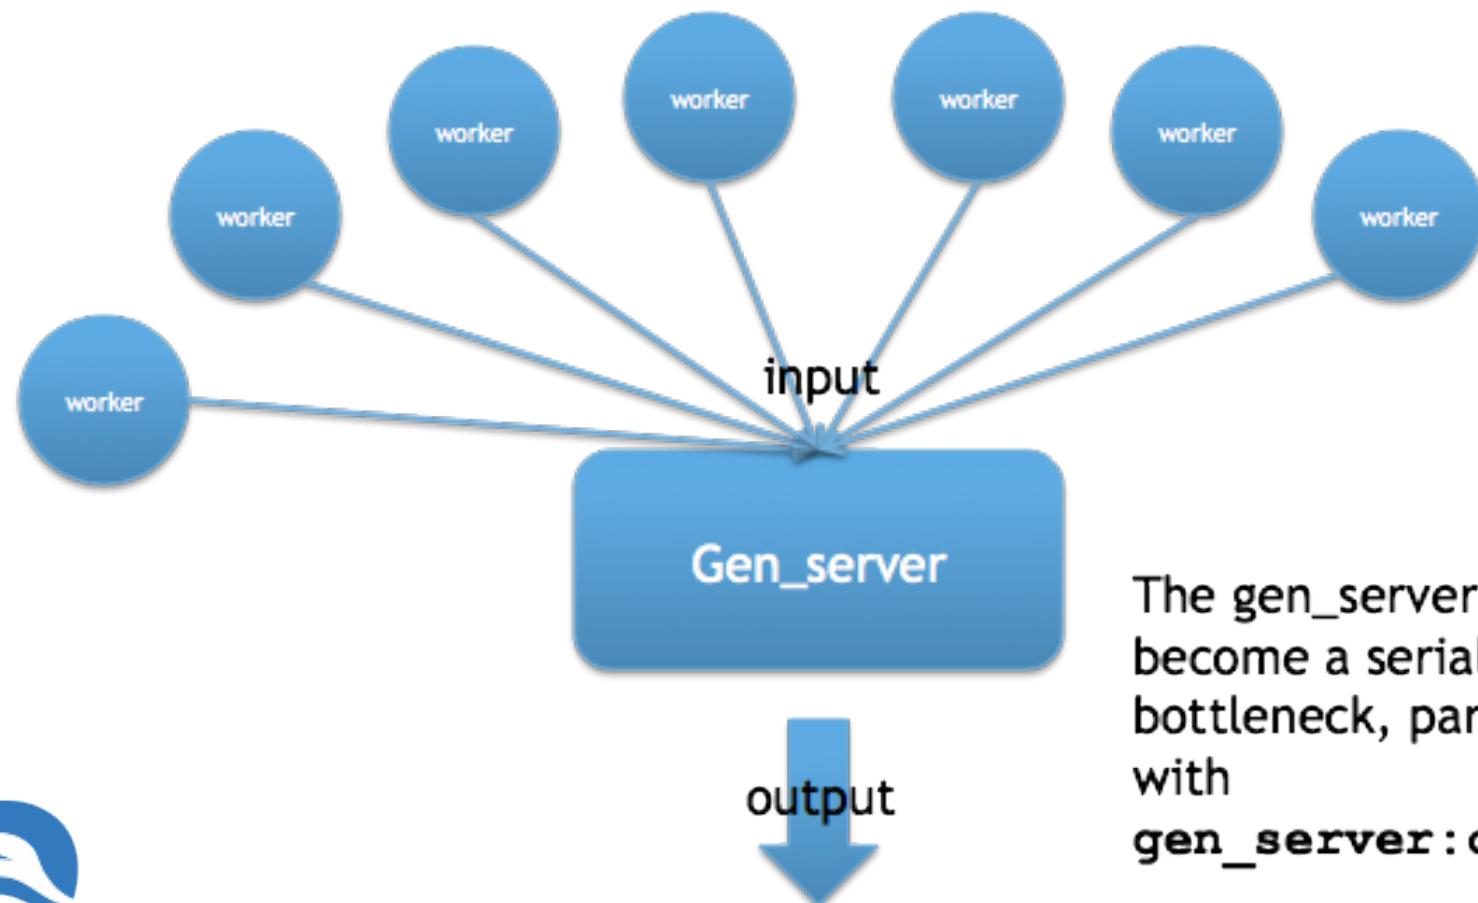
# Now for the Bottlenecks



© 2014 - Erlang Solutions Ltd

*Erlang*  
SOLUTIONS

# Now for the Bottlenecks

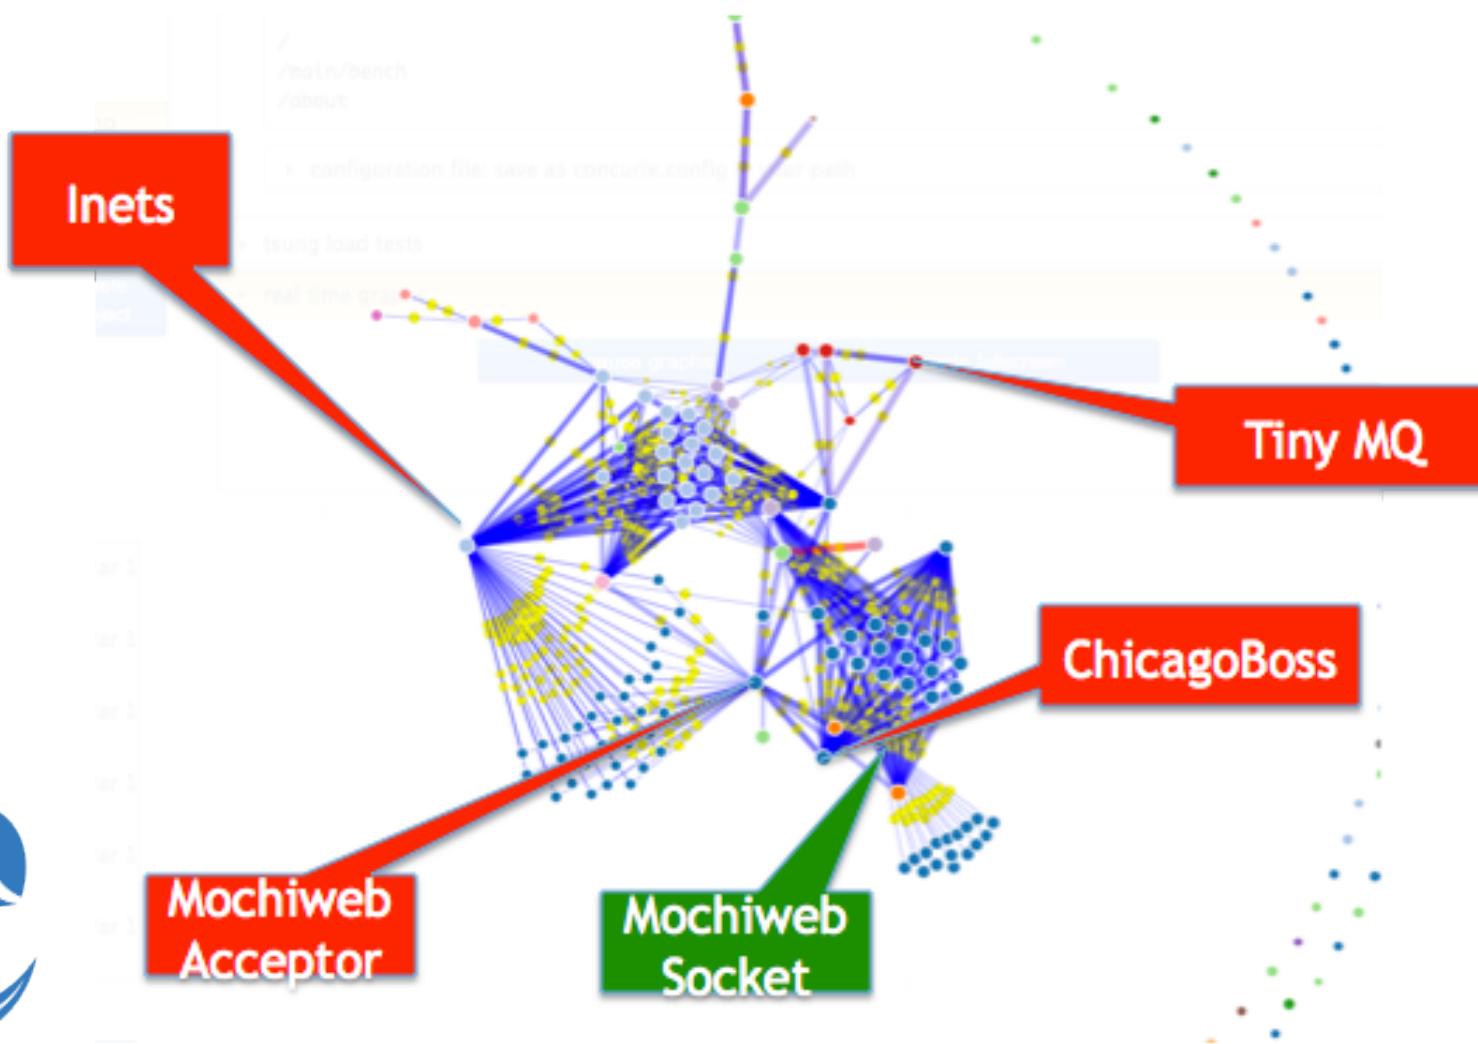


The `gen_server` can become a serialization bottleneck, particularly with  
`gen_server:call(...)`



# Now for the Bottlenecks

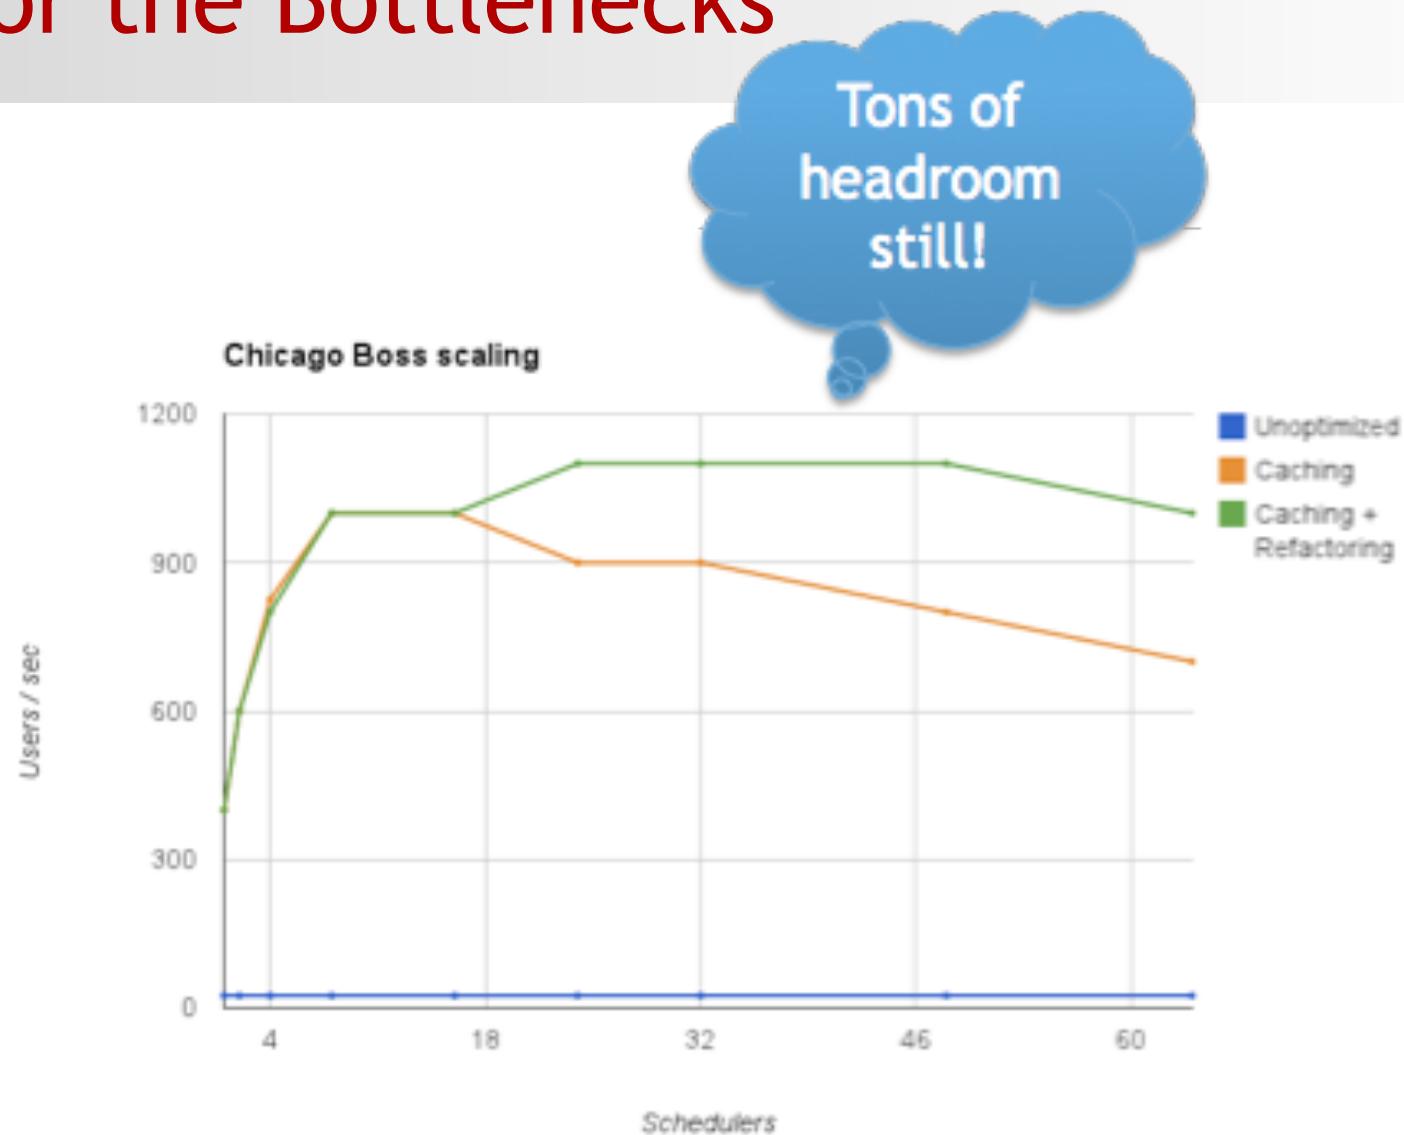
[www.concurix.com](http://www.concurix.com)



© 2014 - Erlang Solutions Ltd

*Erlang*  
SOLUTIONS

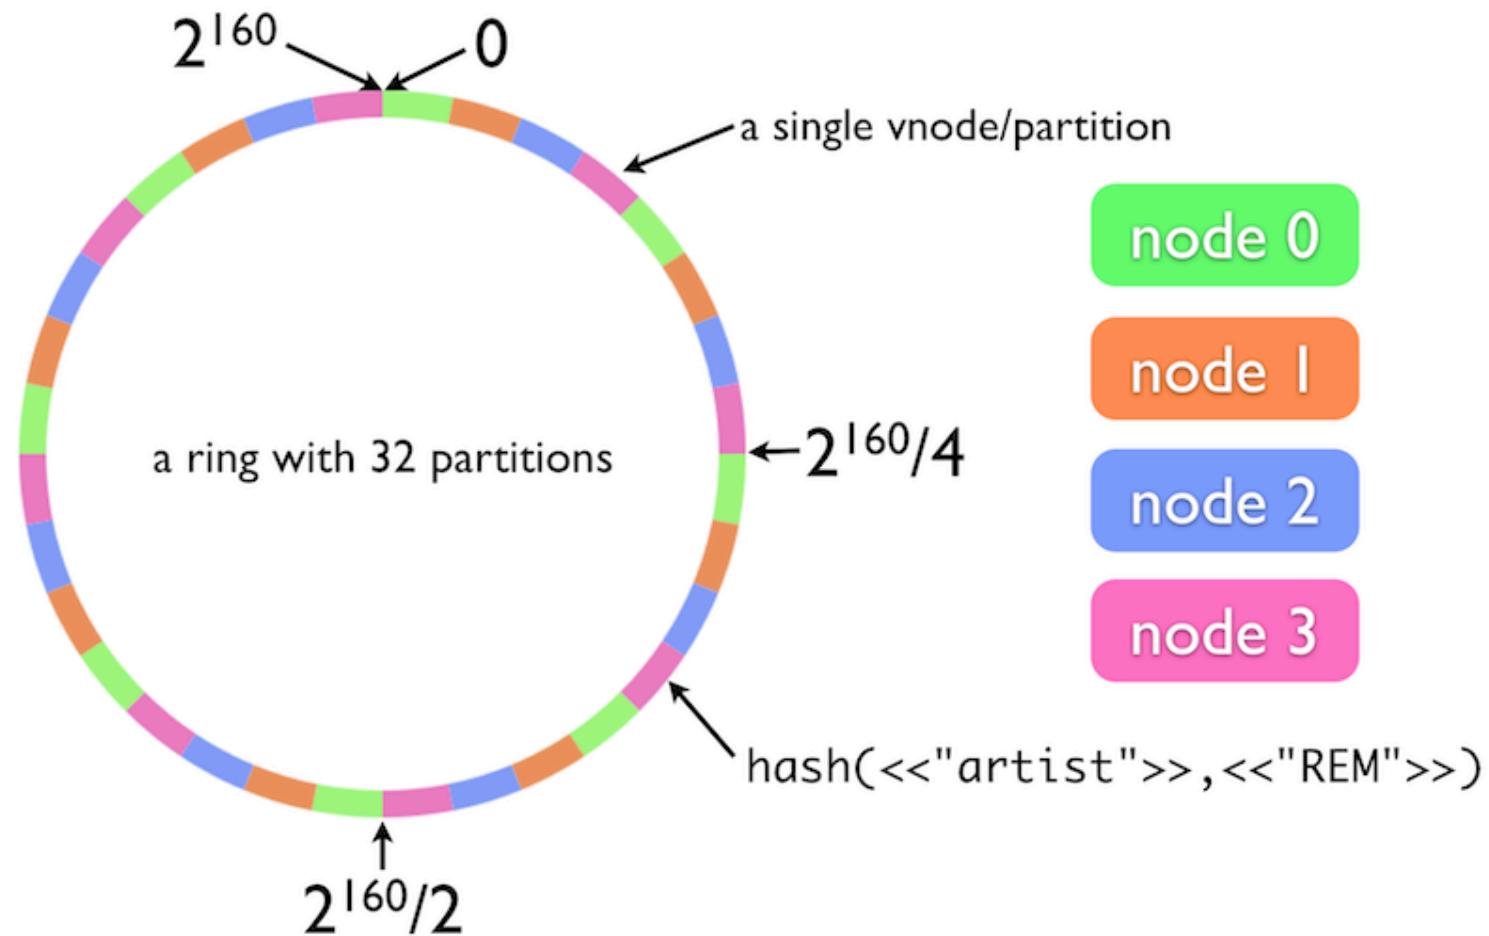
# Now for the Bottlenecks



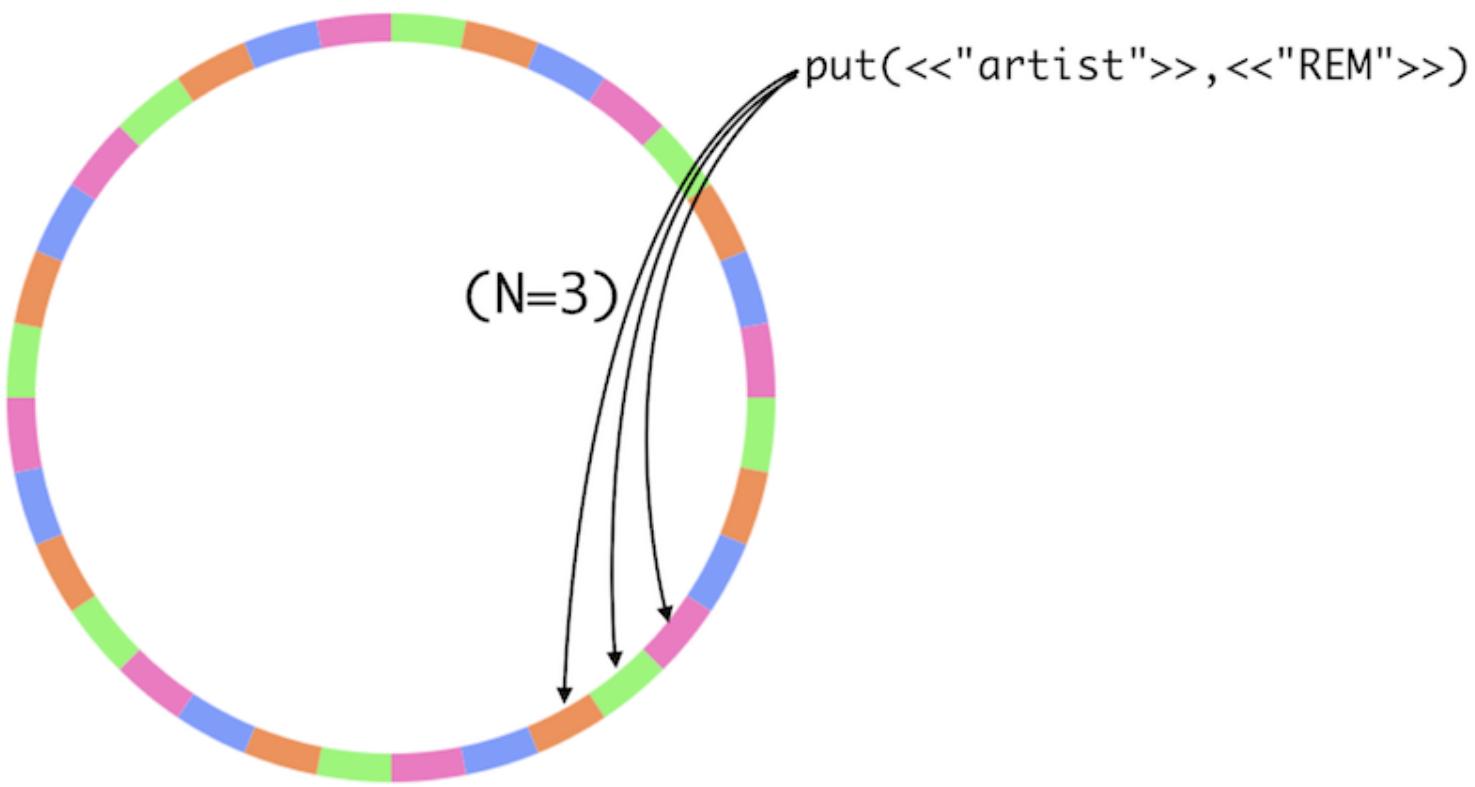
© 2014 - Erlang Solutions Ltd

*Erlang*  
SOLUTIONS

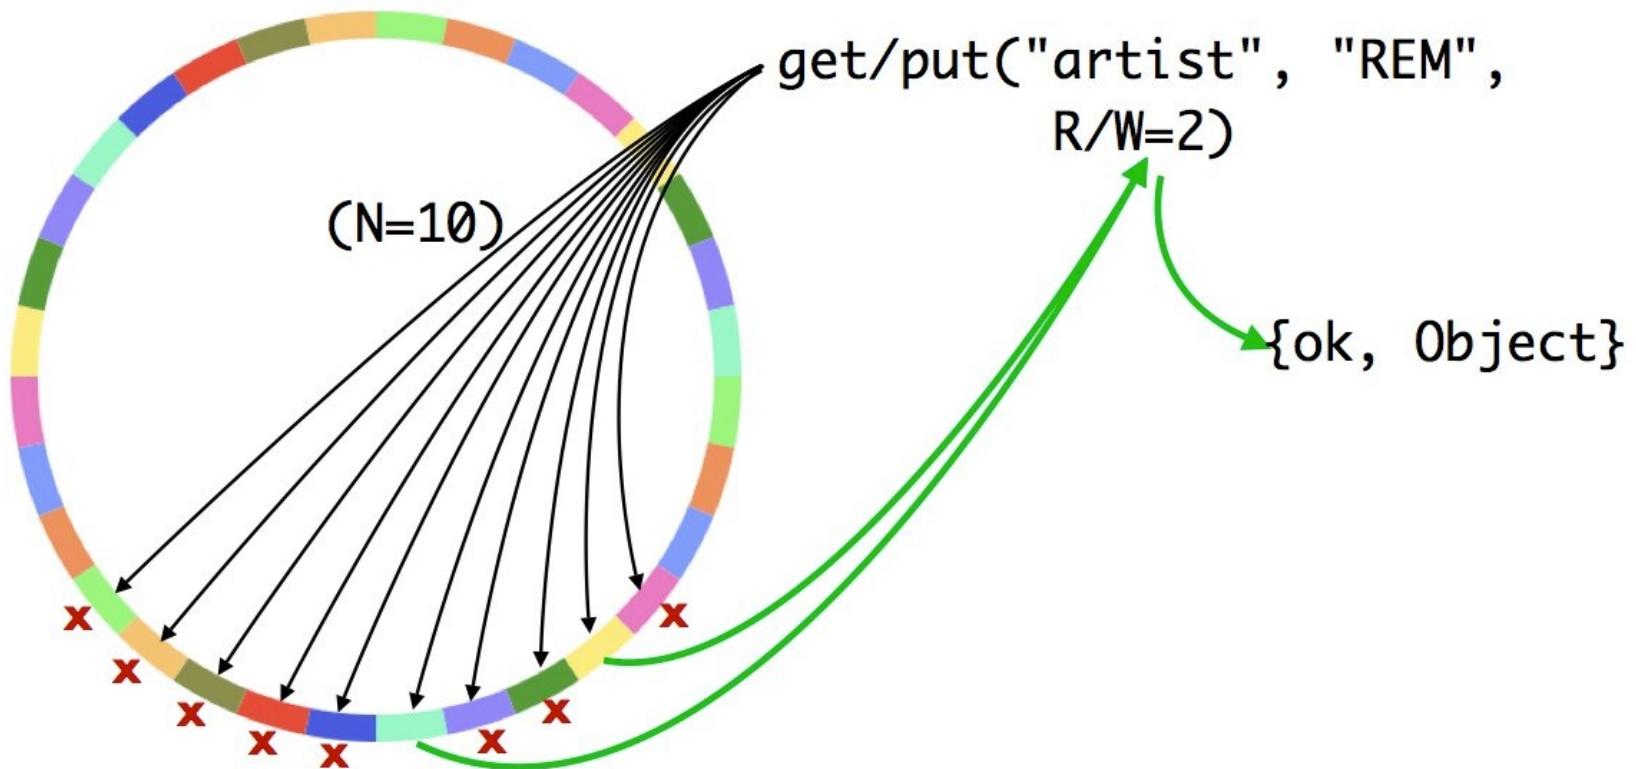
# Riak and other scalable architectures



## N/R/W Values

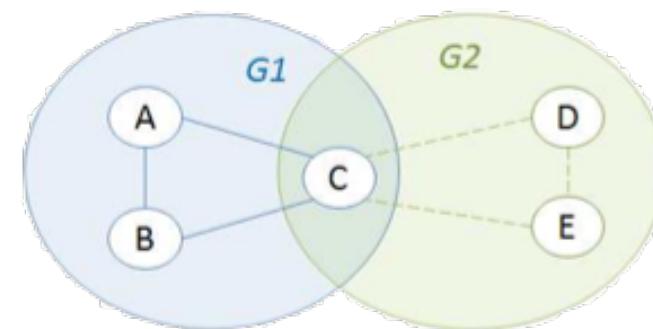
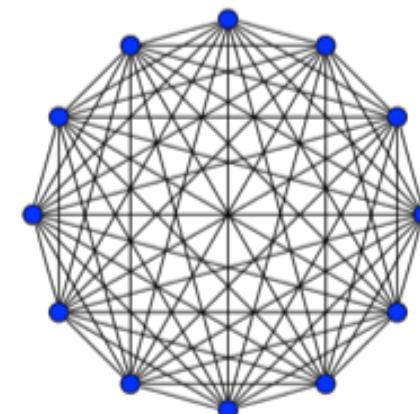


## N/R/W Values



# Clusters and SD Erlang

- TWO MAJOR ISSUES
  - FULLY CONNECTED CLUSTERS
  - EXPLICIT PROCESS PLACEMENT
- SCALABLE DISTRIBUTED (SD) ERLANG
  - NODES GROUPING
  - NON-TRANSITIVE CONNECTIONS
  - IMPLICIT PROCESS PLACEMENT
  - PART OF THE STANDARD ERLANG/OTP PACKAGE
- NEW CONCEPTS INTRODUCED
  - LOCALITY, AFFINITY AND DISTANCE



# Release Statement of Aims



*“To scale the radical **concurrency-oriented programming** paradigm to build **reliable general-purpose software**, such as server-based systems, on **massively parallel machines** ( $10^5$  cores).”*

 RELEASE

University of  
**Kent**



UPPSALA  
UNIVERSITET

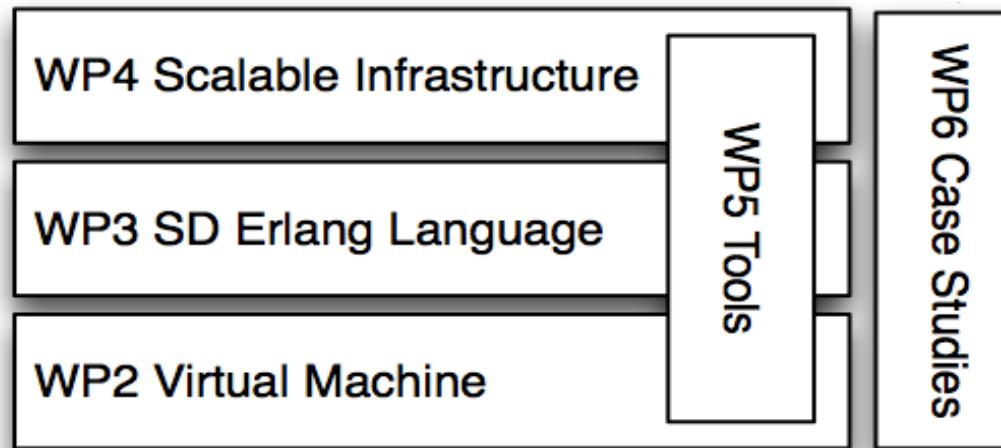


# Release



*“Limitations exist on all levels. You would not want an Erlang VM to run with  $10^5$  schedulers.”*

 RELEASE

The RELEASE icon consists of a grid of nine orange squares arranged in three rows of three.

# Release



Push the responsibility for scalability from the programmer to the VM

Analyze performance and scalability

Identify bottlenecks and prioritize changes and extensions

Tackle well-known scalability issues

Ets tables (shared global data structure)

Message passing, copying and frequently communicating processes

# Thank You!

@francescoc  
francesco@erlang-solutions.com

O'REILLY®

Designing for  
Scalability with  
Erlang/OTP

IMPLEMENTING ROBUST,  
FAULT-TOLERANT SYSTEMS

Early Release  
RAW & UNEDITED



Francesco Cesarini & Steve Vinoski